



PRACTITIONER

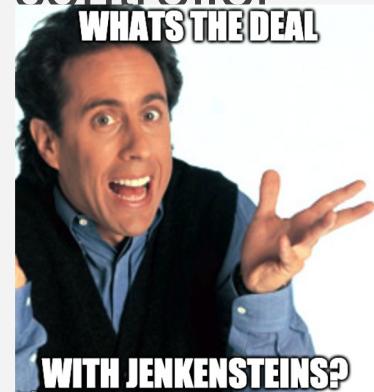
From Big and Slow to Small and Agile: Splitting Monolithic Jenkins Controllers for Increased Performance

Dylan Dewhurst

Sr. Developer Support Engineer at
CloudBees

What's the Deal with Monolithic Jenkins Controllers?

- What is it? The overloaded, overcrowded Jenkins controller that many teams use for all their jobs
- Fairly common occurrence at large enterprise companies
- Easier for Admins to onboard onto existing controller
- This creates problems down the road



PRACTITIONER

Monolith Vs. Horizontal Scaling

Why Monolithic is A Bad Practice

- The main issue is poor performance
 - Slow startup, slow UI, long running builds
- This often leads to vertical scaling:
 - More CPU
 - More memory
 - More disk space
- Outages affect many users and down time is usually increased



Why Monolithic is A Bad Practice

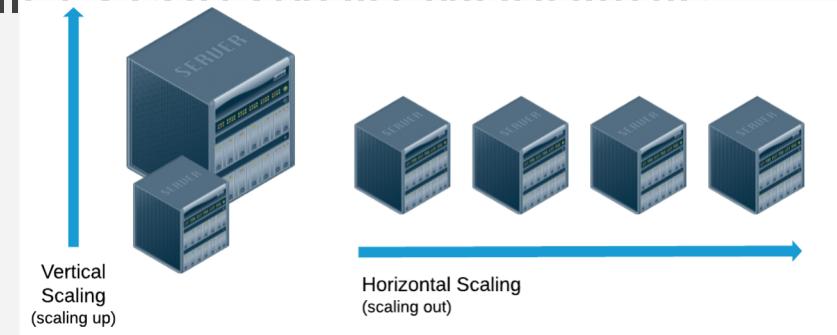
- Jobs affecting performance are harder to find
 - Too many builds running simultaneously
 - Many nested folders of jobs
- Too many plugins. Many are out of date.
 - Sometimes plugins haven't been updated in years
 - There may be plugins not even used
- Managing permissions becomes harder



Horizontal Scaling is the Answer!

- Many smaller Jenkins controllers can avoid Monolithic problems

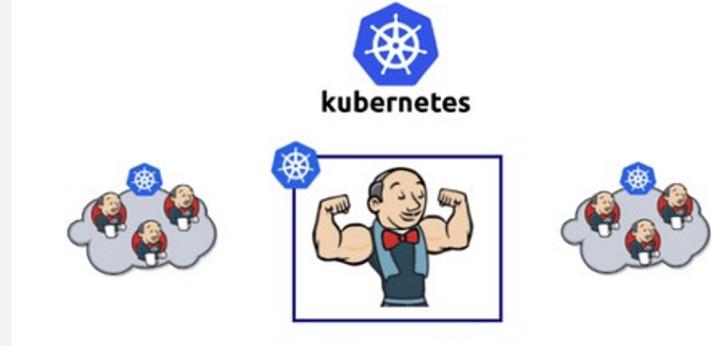
- Faster startup
 - Faster build times
 - Faster UI



- Smaller number of users, typically one team
 - Could self manage
- Less plugins to maintain

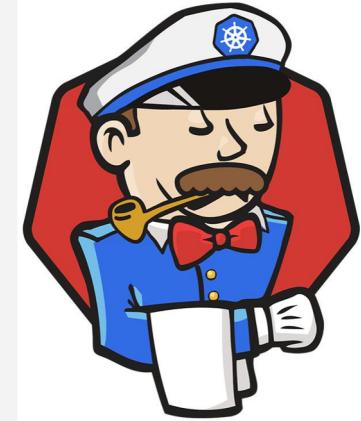
Horizontal Scaling is the Answer!

- An outage is shorter and affects less users
- Easier to find bad plugin or job
- Great for Kubernetes!
 - For CloudBees CI, easily managed by Operations Center
- This all leads to happier users and Admins!



Why Kubernetes?

- Containers work great with single applications
 - Fast to start
 - Easy to move, change config and resources
 - Many fit on one server, easily scalable
- Cloud ready with EKS, GKE, AKS
- Ephemeral Agents that can scale on demand
- Kubernetes is the future of DevOps!



PRACTITIONER

Guidelines for Scaling

Scaling Guidelines - Sizing

- Keep user count low per controller
 - Preferably one team or business unit
- Limit heap size to 16GB maximum
- As best practice, no more than 5000 jobs per controller
- No executors on controller
 - Use Agents for all builds



Scaling Guidelines - Configuration

- Use Configuration as Code plugin
 - For CloudBees CI, Configuration as Code for Controllers
- Limit use of Groovy in Pipeline jobs and Script Console
 - Could heavily impact performance
- Set build discarding policies
- Use webhooks over polling

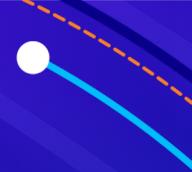


PRACTITIONER

Preparing to Horizontally Scale

Preparing Existing Controller

- Data cleanup
 - Configure build discard policy. Remove unused jobs.
 - Remove unused plugins. If a CloudBees CI customer, can use CloudBees Plugin Usage Analyzer plugin.
- Update in use plugins
- Review Jenkins health using Health Advisor by CloudBees plugin
 - Prevent duplicating issues when scaling



Creating New Jenkins Controller

- Create on a server or container with less resources
 - May require a bit of trial and error to get right
- Use same Jenkins version and same/similar arguments, variables
 - If using containers, most configuration can be reused
- Start the new controller for the first time for initial setup

PRACTITIONER

Data Migration and Configuration

Backup and Export

- Put Source Jenkins in Quiet Down mode: `$JENKINS_URL/quietDown`
- Export system and folder level credentials
 - Scripts provided in CloudBees “[Splitting a controller](#)” documentation
- Backup `$JENKINS_HOME` directory
 - If a CloudBees CI customer, can use the CloudBees Backup plugin
 - Make sure to only backup the jobs in `$JENKINS_HOME/jobs` that are needed for the new controller.

Backup - Excludes

- Exclude the following files from your backup as they are unique to each Jenkins controller:

```
secret.key  
license.xml  
identity.key.enc  
jgroups/  
operations-center-cloud*  
operations-center-client*  
com.cloudbees.opscenter.client.plugin.OperationsCenterRootAction.xml  
nodes/  
secrets/master.key
```

Restoring at Destination

- Copy the backup data into the `$JENKINS_HOME` of the new destination controller
 - If a CloudBees CI customer, can use a restore job
- Import the system and folder level credentials from source controller
 - Use encoded output from export script
 - Scripts provided in CloudBees “[Splitting a controller](#)” documentation



Additional Configuration

- After restore, a number of items need to be reconfigured:
 - Agents
 - Authentication (possibly)
 - Webhooks
 - Jenkins URL
 - Third Party Integrations
 - Another plugin review

Resources

- More detailed steps covered in: [CloudBees Splitting a Controller](#)
- [Configuration as Code plugin](#)
- [Health Advisor by CloudBees](#)
- [Quiet Down mode](#)
- [Kubernetes Plugin](#)
- [Pipeline Best Practices](#)
- [Best Strategy for Disk Space Management](#)
- [SCM Best Practices](#)



Thank you!

Q&A

Dylan Dewhurst
Sr. Developer Support Engineer